

# *Introduction to Text Analysis*

# It may seem obvious but.. what is text?

- From a sociolinguistic perspective: any symbolically encoded language.

These come in many types:

- Alphabetic – symbols (letters) to represent vowels and consonants (e.g., Latin)
- Abjads – symbols to represent consonants, with diacritics (or reused consonants) to represent vowels (e.g., Arabic, Hebrew)
- Syllabic – Symbol systems representing consonants plus inherent vowels (e.g., Devanagari)
- Semanto-phonetic – Symbols that carry meaning and/or sounds (e.g., Chinese)
- Plus, there are additional writing systems: Braille, shorthand, phonetic alphabets

# What is Text?

- From a computer perspective a binary encoding system for storing written language symbols from the world's various systems

Binary	Oct	Dec	Hex	Glyph		
				1963	1965	1967
010 0000	040	32	20			space
010 0001	041	33	21			!
010 0010	042	34	22			"
010 0011	043	35	23			#
010 0100	044	36	24			\$
010 0101	045	37	25			%
010 0110	046	38	26			&
010 0111	047	39	27			'
010 1000	050	40	28			(
010 1001	051	41	29			)
010 1010	052	42	2A			*
010 1011	053	43	2B			+
010 1100	054	44	2C			,
010 1101	055	45	2D			-
010 1110	056	46	2E			.
010 1111	057	47	2F			/
011 0000	060	48	30			0
011 0001	061	49	31			1
011 0010	062	50	32			2

Row	Cells	Range(s)
00	20-7E	Basic Latin (00-7F)
	A0-FF	Latin-1 Supplement (80-FF)
01	00-13, 14-15, 16-2B, 2C-2D, 2E-4D, 4E-4F, 50-7E, 7F	Latin Extended-A (00-7F)
	8F, 92, B7, DE-EF, FA-FF	Latin Extended-B (80-FF ...)
02	18-1B, 1E-1F	Latin Extended-B (... 00-4F)
	59, 7C, 92	IPA Extensions (50-AF)
03	BB-BD, C6, C7, C9, D6, D8-DB, DC, DD, DF, EE	Spacing Modifier Letters (B0-FF)
	74-75, 7A, 7E, 84-8A, 8C, 8E-A1, A3-CE, D7, DA-E1	Greek (70-FF)
04	00-5F, 90-91, 92-C4, C7-C8, CB-CC, D0-EB, EE-F5, F8-F9	Cyrillic (00-FF)
1E	02-03, 0A-0B, 1E-1F, 40-41, 56-57, 60-61, 6A-6B, 80-85, 9B, F2-F3	Latin Extended Additional (00-FF)
1F	00-15, 18-1D, 20-45, 48-4D, 50-57, 59, 5B, 5D, 5F-7D, 80-B4, B6-C4, C6-D3, D6-DB, DD-EF, F2-F4, F6-FE	Greek Extended (00-FF)
20	13-14, 15, 17, 18-19, 1A-1B, 1C-1D, 1E, 20-22, 26, 30, 32-33, 39-3A, 3C, 3E, 44, 4A	General Punctuation (00-6F)
	7F, 82	Superscripts and Subscripts (70-9F)
	A3-A4, A7, AC, AF	Currency Symbols (A0-CF)

From ASCII (7-bits) to Unicode: An international standard that supports up to four bytes (32 bits) per symbol and that encodes 137,439 characters from 146 different scripts.


# What is text?

- From a data scientist's perspective, a set of documents organized into a corpus, where each document contains an encoded sample of written language, generally:
  - Human readable files (e.g., plain text)
  - Metadata describing data source, date of capture, speaker/author, etc.
- Character/symbol encoding appropriate to language (Arabic, Hebrew..)
- One file per document, or container file includes document separator characters (e.g., newline)
- Each file contains unstructured, “natural” language content – generally the same “type” of content for each file in the corpus (i.e., don't mix tweets with book chapters)



# “Unstructured” is the Key

- Natural language is notoriously flexible and ambiguous
  - even simple tasks like parsing are complicated:
- Contractions; Compound words; Misspellings; Proper Names; Preposition Attachment; Vernacular
- Transforming a set of documents into data that allow for comparisons, linkages, visualization, or other analysis, is complex and contains many decisions



# Analytical Overview

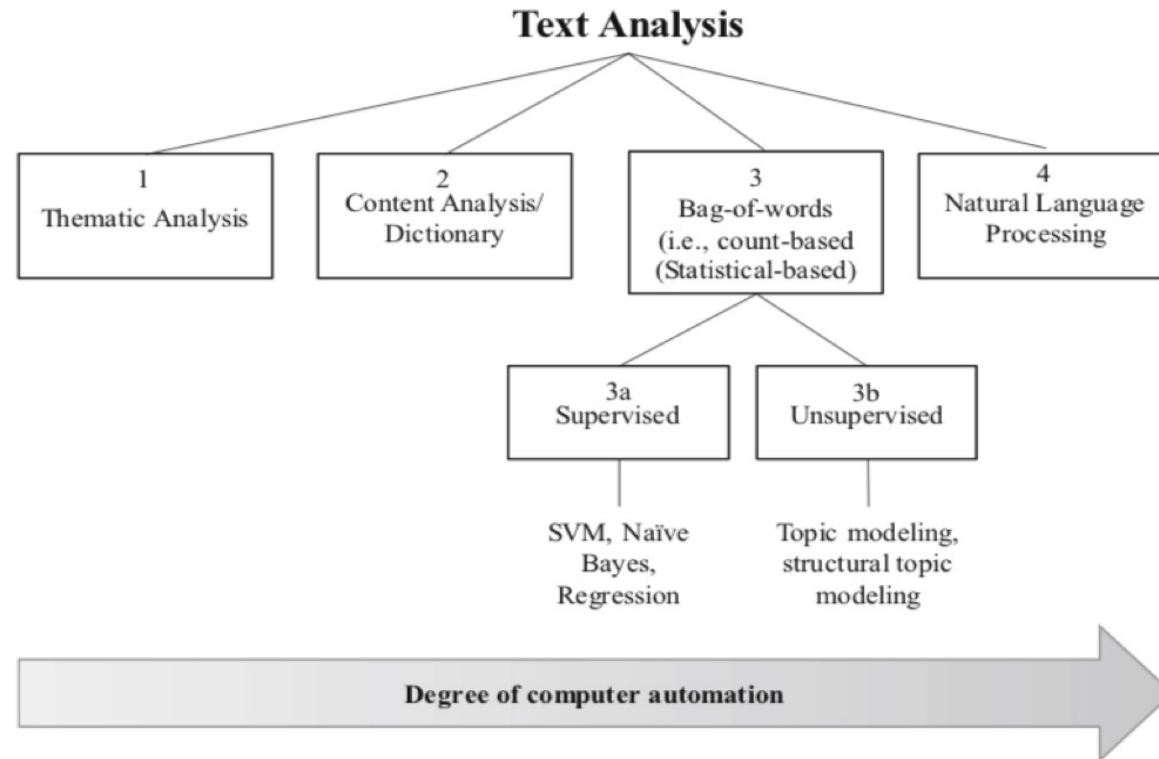
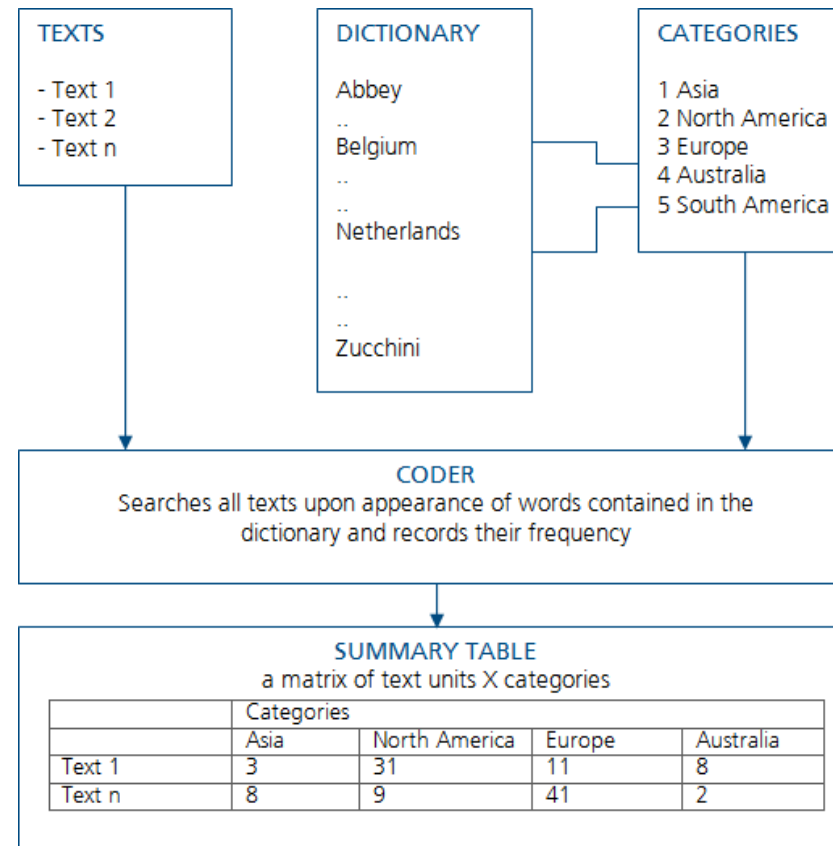


Image Credit: Banks et al. (2018), Fig. 1, p. 447

# Dictionary-Based Content Analysis

- Simplest form of automated text analysis
- Reduce texts to frequency counts of various vocabulary words
- Lookup vocabulary words in appropriate dictionaries – e.g., positive/negative sentiment
- Quantify results based on connection between word frequency and dictionary connotations
- Results can be used to compare documents to one another or to summarize corpora



# Statistical Approaches

	I	love	dogs	hate	and	knitting	is	my	hobby	passion
Doc 1	1	1	1							
Doc 2	1		1	1	1	1				
Doc 3					1	1	1	2	1	1

- Treat each text fragment as a collection of units/words, without regard to word order
- Use many text fragments – e.g. chapters from a book; each fragment is considered a “document”
- Compute **a term-document matrix (TDM)** (or document term matrix); a sparse matrix pinning down the count of appearance of a term (aka word) in a document
- Conduct statistical or machine learning analysis on the TDM to reveal patterns



## Example TDM

Documents

However, complexity  
We will see how small  
Given a function-based  
Using entropy of traffic  
We study the complexity  
of influencing elections  
through bribery: How  
computationally complex  
is it for an external actor  
to determine whether by  
a certain amount of  
bribing voters a specified  
candidate can be made  
the election's winner? We  
study this problem for  
election systems as varied  
as scoring ...

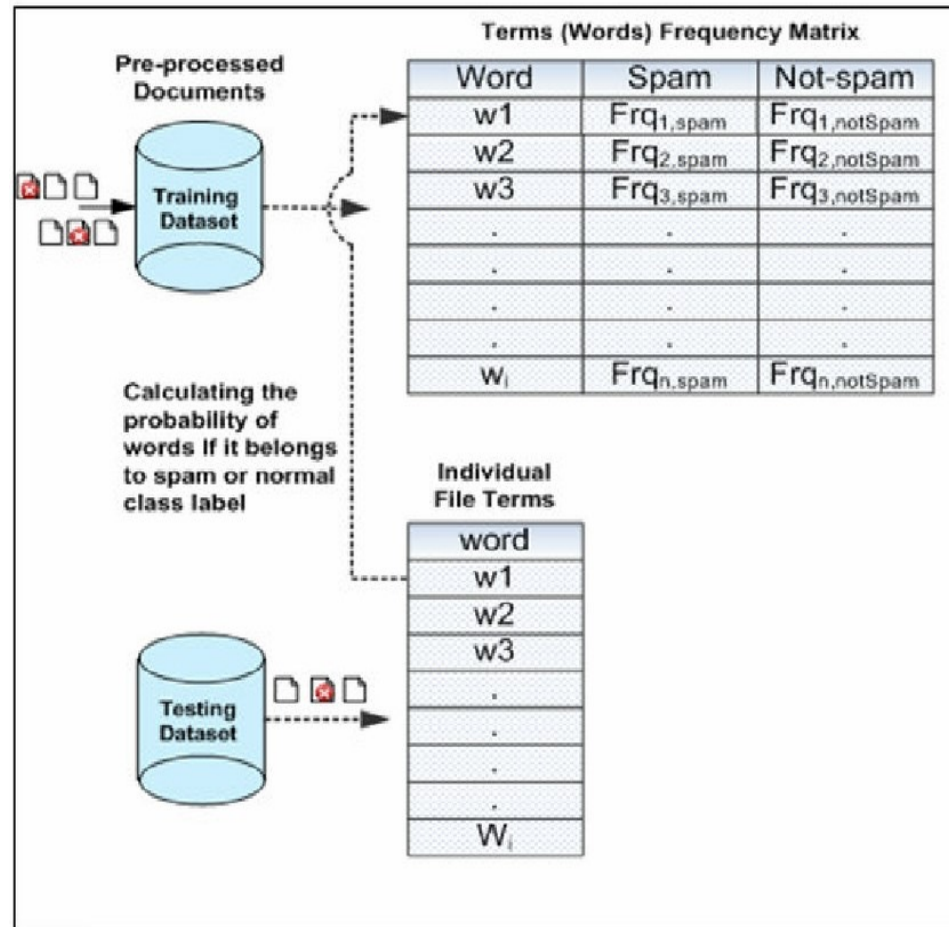


Vector-space  
representation

	D1	D2	D3	D4	D5
complexity	2		3	2	3
algorithm	3			4	4
entropy	1			2	
traffic		2	3		
network		1	4		

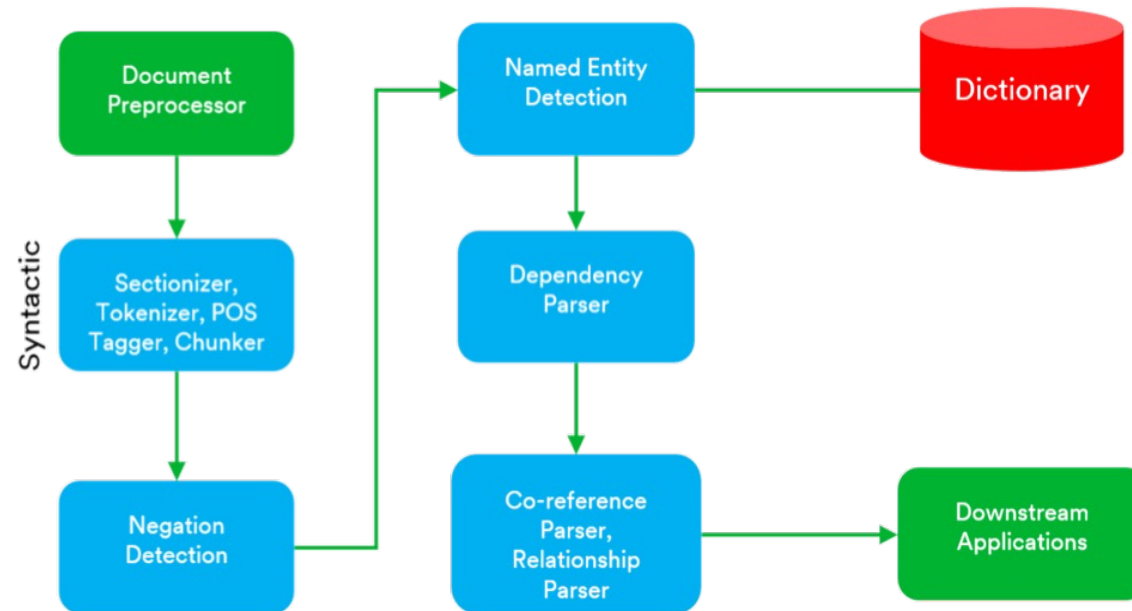
Term-document matrix

# Example Application: Spam Filter



# Natural Language Processing

- Tokenizes and tags text to detect and preserve **part-of-speech**
- Post processing of tokens tries to resolve syntactic dependencies, co-references, names of entities, sequences of events, etc.
- Machine learning can then be used to solve a variety of difficult challenges such question answering, information retrieval, and machine translation, e.g., Alexa, Siri...



## **Research Questions to Address with Text**

- Organizations: Examine the dynamics of team formation by analyzing listserv emails
- Economics: Analyze the content of job postings to estimate the demand for particular jobs across years/regions
- Sociology: Track levels of hate speech over time using text from online forums
- Politics: Detect the proportion of fear-based appeals in campaign rhetoric based on party affiliation
- Psychology: Conduct sentiment analysis on open ended comments in surveys to detect faking

# Text Data Sources

- <https://github.com/niderhoff/nlp-datasets>

Alphabetic list of freely available text corpora

- <https://www.gutenberg.org>

Project Gutenberg, more than 57,000 public domain books available in various formats

- [https://en.wikipedia.org/wiki/List\\_of\\_text\\_corpora](https://en.wikipedia.org/wiki/List_of_text_corpora)

List of corpora for more than 30 languages including “parallel corpora” where the same text is available in two or more languages

- <https://gengo.ai/datasets/the-best-25-datasets-for-natural-language-processing/> As the URL suggests. . .

## **Worked Example 1:**

Importing Text and Dictionary Analysis

# Overview

- Read in a text file, separating into paragraphs, each of which will be treated as a separate “document”
- Tokenize the text into a bag of words, representing it as a Document-Term Matrix
- Extract term frequency information
- Visualize term frequency information
- Match term frequency to dictionaries of positive and negative connoted words (merge)
- Visualize results
- Research question: In a nomination acceptance speech, do positive sentiments predominate over negative ones?

# Read in a Speech

```
library(tm)
charVector <- scan("speech.txt", character(0), sep = "\n")
head(charVector)
```

```
1"I speak tonight of gratitude, achievement, and high hopes
for our country."
2 "Tonight, I think first of those who helped get me here
- starting with the people of Tennessee. Then, those who
braved the first snows of Iowa and New Hampshire -- and all of
you here, from all over this country, who have come with me
into the warm sunlight of this great city."
3 "While I can't thank each of you individually in words,
I do so in my heart."
4 "And I know you won't mind if I single out someone who
has just spoken so eloquently, someone I've loved with my whole
heart since the night of my high school senior prom -- my wife,
Tipper. We've been lucky enough to find each other all over
again at each new stage of our lives - and we just celebrated
our 30th wedding anniversary."
```



# Read in Dictionaries

```
posWords <- scan("positive-words.txt", character(0), sep = "\n") #  
2006 items  
negWords <- scan("negative-words.txt", character(0), sep = "\n") #  
4783 items  
head(posWords,15)  
head(negWords,15)
```

```
> head(posWords,15)  
[1] "a+"          "abound"      "abounds"     [4]  
"abundance"    "abundant"    "accessable"  
[7] "accessible"  "acclaim"     "acclaimed"  
[10] "acclamation" "accolade"    "accolades"   [13]  
"accommodative" "accomodative" "accomplish"  
> head(negWords,15)  
[1] "2-faced"      "2-faces"     "abnormal"    "abolish"  
[5] "abominable"   "abominably"  "abominate"   "abomination"  
[9] "abort"        "aborted"     "aborts"      "abrade"  
[13] "abrasive"     "abrupt"      "abruptly"
```

# Bag of words: Two class transformations

```
wordVector <- VectorSource(charVector)
wordCorpus <- Corpus(wordVector)
class(wordVector); typeof(wordVector); length(wordVector)
class(wordCorpus); typeof(wordCorpus); length(wordCorpus)
```

```
> class(wordVector); typeof(wordVector); length(wordVector)
[1] "VectorSource" "SimpleSource" "Source"
[1] "list"
[1] 166
> class(wordCorpus); typeof(wordCorpus); length(wordCorpus)
[1] "SimpleCorpus" "Corpus"
[1] "list"
[1] 166
```

# Token Clean Up


```
# first step transformation: make all of the letters in "wordCorpus"
lowercase wordCorpus <- tm_map(wordCorpus, content_transformer(tolower))

# second step transformation:

#remove the punctuation in "wordCorpus" wordCorpus <- tm_map(wordCorpus,
removePunctuation)

# third step transformation: remove numbers in "wordCorpus" wordCorpus <-
tm_map(wordCorpus, removeNumbers)

# final step transformation: take out the "stop" words, such as "the", "a"
and "at" wordCorpus <- tm_map(wordCorpus, removeWords,
stopwords("english")) wordCorpus[["1"]][["content"]] # Review what's left
of the first paragraph
```



Ignore warnings  
about dropped  
documents.

```
> wordCorpus[["1"]][["content"]]
[1] " speak tonight  gratitude achievement  high hopes  country"
```

# Create a term-document matrix "tdm"

```
tdm <- TermDocumentMatrix(wordCorpus)
```

```
tdm
```

```
View(tdm)
```

Name	Type	Value
tdm	list [1211 x 166] (S3: TermDo	List of length 6
i	integer [2574]	1 2 3 4 5 6 ...
j	integer [2574]	1 1 1 1 1 1 ...
v	double [2574]	1 1 1 1 1 1 ...
nrow	integer [1]	1211
ncol	integer [1]	166
dimnames	list [2]	List of length 2

```
> tdm
<<TermDocumentMatrix (terms: 1211, documents: 166)>>
Non-/sparse entries: 2574/198452
Sparsity           : 99%
Maximal term length: 18
Weighting          : term frequency (tf)
```

# Some things to do with a TDM

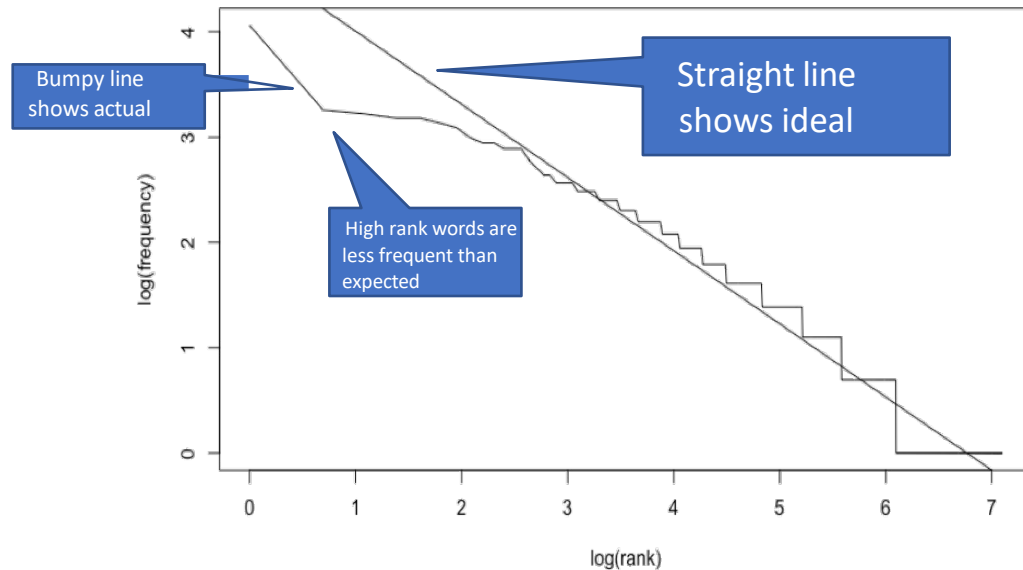
```
fTerms <- findFreqTerms(tdm, lowfreq = 20)
# Common terms: "tonight" "new" "people"
# "children" "family" "will" "families"
"america" findAssocs(tdm, fTerms, 0.4)
```

Note: Excerpted results.

```
> findAssocs(tdm, fTerms, 0.4)
$people fighting      champion      charged constitution      district
      0.62      0.53      0.53      0.53      0.53
$children raising
      0.4
$families tax working
      0.41      0.40
```

# Zipf's Law: Zipf\_plot(tdm)

- The numeric frequency of any word in an NL corpus is negatively related to its frequency rank.
- Example: Most frequent word occurs 2X as much as the next most frequent word, 3X as much as the third most frequent word, etc. . .



# Aggregate TDM into vector of word counts

```
m <- as.matrix(tdm) # Coerce to regular matrix

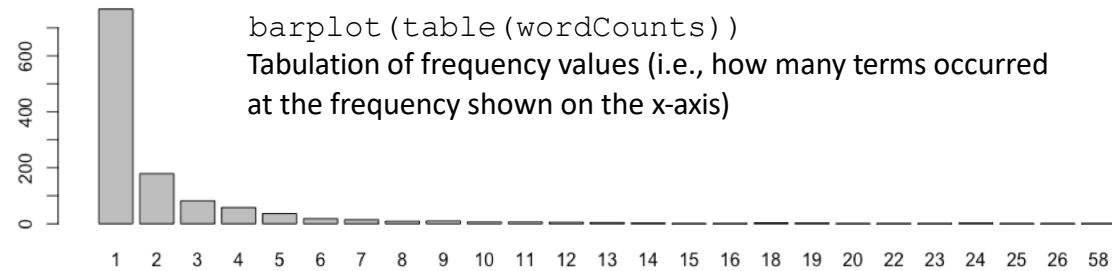
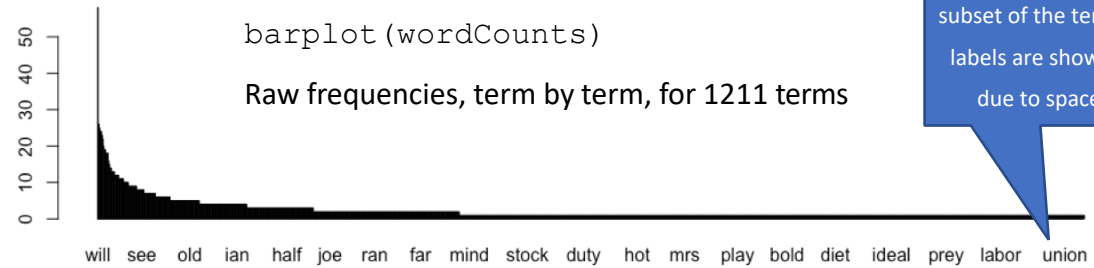
# create a list of counts for each word named "wordCounts"
wordCounts <- rowSums(m)

# sort words in "wordCounts" by frequency wordCounts
<- sort(wordCounts, decreasing=TRUE)

# check the first several items in "wordCounts" to see if it is
built correctly head(wordCounts)
totalWords <- sum(wordCounts) # Calculate the total number of occurrences
```

```
> head(wordCounts)
will families    people    new america  family
   58      26      25     24     24     23
```


## Word Count Bar Plots






# Positive Word Matches: 7.8% of total

```
matchedP <- match(names(wordCounts), posWords, nomatch = 0)
matchedP <- matchedP != 0 # Make a Boolean map
matchedP <- wordCounts[matchedP] # Apply Boolean map
barplot(matchedP, las=2, cex.names=0.75) # Small vertical labels
sum(matchedP)/totalWords # Proportion of pos words vs total
```



# Exercise 14.1: Generate the Negative Word Matches and Interpret Results



## **Answering a Research Question**

- Dictionary analysis indicated that positive terms occurred twice as frequently as negative terms in this speech
- This result affirms that idea that a nomination acceptance speech is typically an affirmative celebration of success in the primary
- Other recent nomination speeches should be similarly analyzed to see if the pattern holds

### **Exercise: Analyze another speech**

Analyze another presidential acceptance speech using the same code

- Go to:

<https://www.presidency.ucsb.edu/documents/appcategories/elections-and-transitions/conventionspeeches>

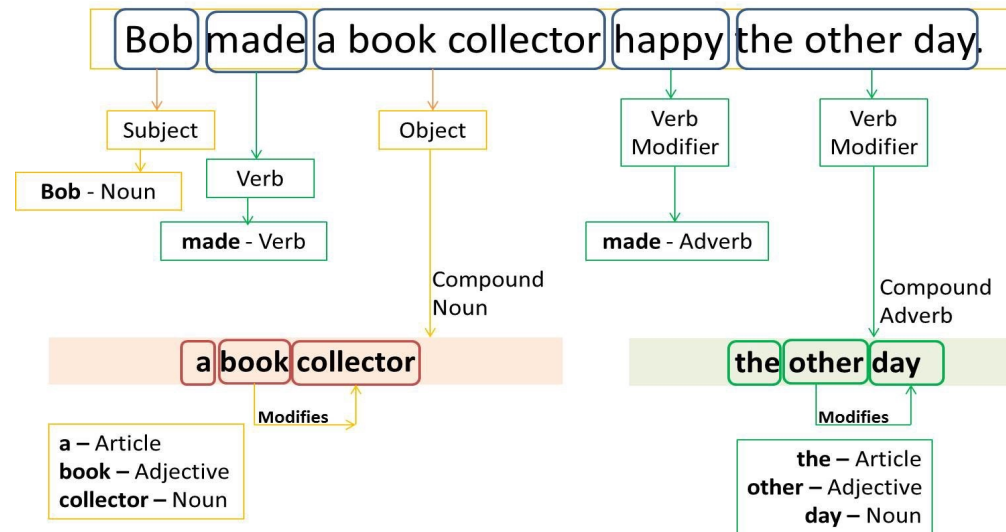
to obtain another speech

- Copy and paste the nominee's remarks from the web page into a plain text file on your computer
- Check to make sure that there are line breaks for each paragraph
- Use the same positive and negative word dictionaries

# Brief Review of POS Tagging

# Part of Speech Tagging

- NLP parsers contain extensive, language-specific logic for decomposing sentences (tokenizing) into parts of speech



Prepositions – in, on, at, with  
Conjunctions – but, or, and, for

# Example using spacyr

- spacyr is an R package “wrapper” around an open source Python package called spaCy
- v2.0 of spaCy is one of the most accurate and fastest parsers available

```
spacy_initialize() # Start the spacy session

inText <- c(doc1="Thanksgiving was always special at my
house. Sarah pointed to my sisters and said, 'They are
eating apples.'",
           doc2="Then Sarah pointed to the bowl that
held the fruit and said, 'They are eating apples.'")
```

# Tokenized Output

	doc_id	sentence_id	token_id	token	lemma	pos	entity
1	doc1	1	1	Thanksgiving	thanksgiving	PROPN	DATE_B
2	doc1	1	2	was	be	VERB	
3	doc1	1	3	always	always	ADV	
4	doc1	1	4	special	special	ADJ	
5	doc1	1	5	at	at	ADP	
6	doc1	1	6	my	PRON-	ADJ	
7	doc1	1	7	house	house	NOUN	
8	doc1	1	8	.	.	PUNCT	

Each token uniquely  
identified by its  
position in a sentence  
and document.

Lemmatization resolves  
various forms of each  
word to its core  
meaning/function.

Each token belongs to a  
particular token  
category such as  
"Proper Noun."



# Entity Detection

- Parsers usually contain dictionaries and rules that support detection of named entities such as people, places, and organizations

```
> outTokens2 <- spacy_parse(inText, lemma = FALSE)
> entity_extract(outTokens2)
  doc_id sentence_id entity entity_type
1   doc1           2  Sarah    PERSON
2   doc2           1  Sarah    PERSON
```

# Dependency Tagging

- Syntactical dependencies show, for example, which noun an adjective modifies

	doc_id	sentence_id	token_id	token	head_token_id	dep_rel	entity
9	doc1	2	1	Sarah	2	nsubj	PERSON_B
10	doc1	2	2	pointed	2	ROOT	
11	doc1	2	3	to	2	prep	
12	doc1	2	4	my	5	poss	
13	doc1	2	5	sisters	3	pobj	
14	doc1	2	6	and	2	cc	
15	doc1	2	7	said	2	conj	

Coordinating conjunction  
"and" and conjunct "said"  
are attached to the root  
action, "pointed."

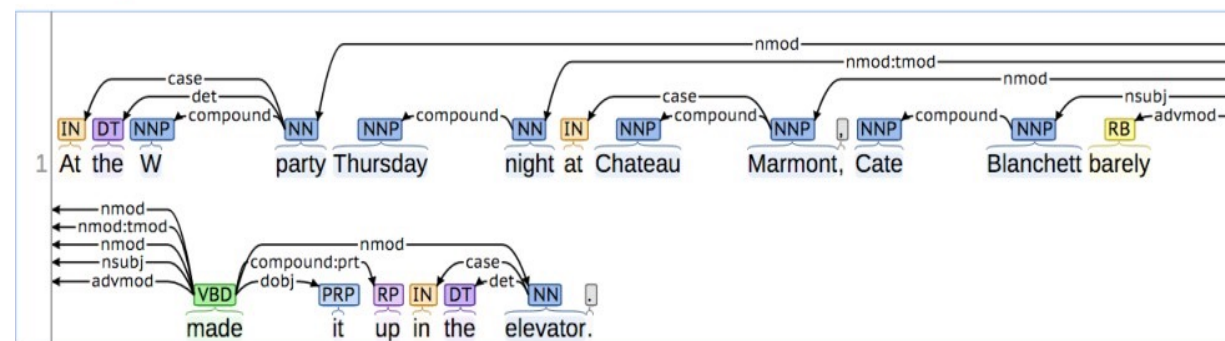
Possessive "my" refers to  
objective of preposition  
"sisters".

# Harder Problem: Detecting Co-Reference

## Named Entity Recognition:

	<div>Date</div>	<div>Time</div>	<div>Location</div>	<div>Person</div>	
1	At	the	W party Thursday night	at Chateau Marmont,	Cate Blanchett barely made it up in the elevator.

## Basic Dependencies:



	<div>Person</div>	<div>Loc</div>	<div>ORDINAL</div>	<div>Location</div>	
1	President Xi Jinping	of China,	on his first	state visit to the United States,	showed off his familiarity with
	<div>Misc</div>		<div>Date</div>	<div>Time</div>	
	American	history and pop culture	on Tuesday	night.	

## Coreference:

	<div>Mention</div>	-----Coref-----	<div>M</div>	
1	President Xi Jinping of China, on his first state visit to the United States,		showed off his familiarity with American history and pop culture on Tuesday night.	

# Major Challenges in NLP: Polysemy and Homonymy

- Polysemy refers to the phenomenon that a word has more than one meaning.

face: the front of the head

a surface of a thing

a person's countenance

a person

"lots of new *faces* around here!"

- Homonymy refers to the phenomenon that two or more words have the same form, but have different meanings.


lie: make an untrue statement.

lie: put oneself in a resting position.

---

# Major Challenges in NLP: Preposition Attachment

I saw the man on the hill with a telescope.



1. I saw the man. The man was on the hill. I was using a telescope.
2. I saw the man. I was on the hill. I was using a telescope.
3. I saw the man. The man was on the hill. The hill had a telescope.
4. I saw the man. I was on the hill. The hill had a telescope.
5. I saw the man. The man was on the hill. I saw him using a telescope.

# Elliptical Construction

- Elliptical construction occurs when words are left out because they are generally understood to be there. Several types:

Type	Example
Verb Gapping	Jeff can play the bass, and Larry . . . the guitar.
Verb Phrase	Jeff wanted to buy the instrument and did . . .
Nominal	Larry played one song, then Jeff played two . . .
	. . . and many more!

- Parsers have difficulty inferring the missing phrases, making later steps (like dependency or co-reference) more difficult/less accurate.

# Philosophical Pitfalls:

## Semantics Versus Pragmatics

- Authentic human communication works in an interpersonal context of shared meanings, personal histories, and prevailing circumstances
- In this perspective, many texts, on their own, do not contain sufficient information to truly convey the underlying meaning of the utterances
- Non-propositional speech acts (when captured in text) lose meaning when separated from their related performances. Example: “I’ll finish the job now.”
- Give an example of an utterance whose meaning varies depending on context. . .

# **Worked Example #2: Topic Modeling of Text**



# Topic Modeling

- Considered an *unsupervised* data/text mining technique: Does a “bottom-up” analysis of documents, topics, and words with no “criterion” (ground truth) to train against
- Extracts topics that are characterized by a combination of more frequent words
- Each document in a corpus comprises a mixture of topics

## Example Using National Anthems

```
> anthems <- read.csv('anthems.csv')
```

```
> charVector <- anthems$Anthem
```

```
> head(charVector, 1)
```

```
[1] "Around our flag we stand united, With one wish  
and one goal, A sacred oath we bestow upon it  
Proclaiming loyalty for our salvation. From war  
abstains only he, Who a traitor is born, He who is a  
true man is not frightened, But dies a warrior to the  
cause. With weapons in our hands a-brandished, We will  
defend our fatherland, Our sacred rights  
we will not relinquish ..."
```

## **library("quanteda")**

- The quanteda package was developed by Kenneth Benoit (<http://kenbenoit.net>) at the London School of Economics
- A comprehensive text processing package that uses some C++ and Fortran – but unlike Stanford Core NLP tools no Java; mostly runs faster and with less memory than Stanford Core NLP
- Includes readtext(), a wrapper for a wide range of text files (PDF, CSV, text, XML, JSON) that simplifies reading in folders of text files
- Directly supports document-level variables for machine learning

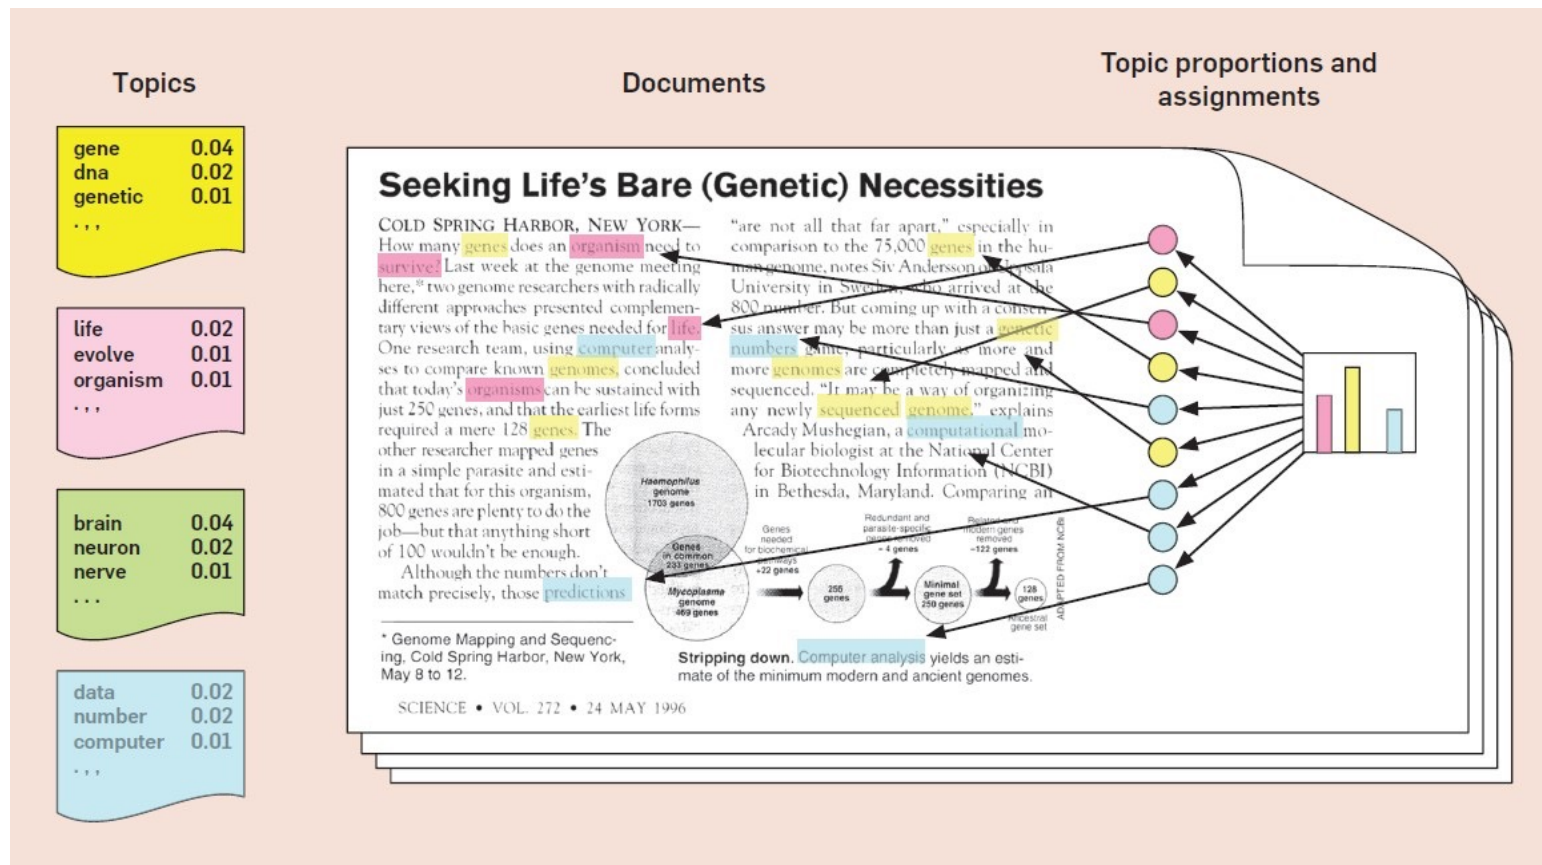
# Preparing and Visualizing Corpus

```
anthemcorpus <- corpus(charVector, docnames=anthems$country)
paras <- corpus_reshape(anthemcorpus, to="paragraphs")
anthem_dtm <- dfm(paras, stem=TRUE, remove_punct=TRUE, remove_symbols =
TRUE, remove_numbers = TRUE, remove=c(stopwords("english")))
anthem_dtm <- dfm_remove(anthem_dtm, c("s","ã","â",'thi'))
anthem_dtm <- dfm_trim(anthem_dtm, min_termfreq=20)
anthem_dtm
colnames(anthem_dtm)
library("quanteda.textplots")
textplot_wordcloud(anthem_dtm) # visualize words
```


```
>textplot wordcloud(anthem dtm)
```



# Using Latent Dirichlet Allocation to Infer Topics from Words and Documents



# Dirichlet Distributions at Different Alpha Levels



```
# 255 docs, where each belongs to  
# one or two topics  
x = rdirichlet(255,  
c(0.1,0.1,0.1)) # Alpha = 0.1  
scatterplot3d(x,scale.y=0.5,angle  
=20,col.axis="red",pch=20)
```

```
# 255 docs, with mixing of 1,2,3  
# topics across docs  
x = rdirichlet(255, c(1,1,1))  
# Alpha = 1  
scatterplot3d(x,scale.y=0.5,angle  
=20,col.axis="red",pch=20)
```

```
# 255 docs, all docs include  
# mix of all three topics  
x = rdirichlet(255, c(10,10,10))  
# Alpha = 10  
scatterplot3d(x,scale.y=0.5,angle  
=20,col.axis="red",pch=20)
```

LDA is a Bayesian technique, guided by two hyperparameters of the Dirichlet distribution, alpha and beta. Low “prior” alpha value has each document composed of only a few key topics. (good separation) Low “prior” beta has each topic composed of only a few key words. (more dense)

# Estimate LDA Model

```
# topic model with 5 total topics (more on how to choose the number of topics)
topic_model <- LDA(anthem_topics, method = "VEM", k=5) # Latent Dirichlet Allocation
termsOut<- terms(topic_model, 8) # Descriptive stems for each topic
```

```
> terms(topic_model,8)
```

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5
[1,]	"us"	"land"	"glori"	"nation"	"may"
[2,]	"peopl"	"o"	"die"	"god"	"homeland"
[3,]	"let"	"love"	"fatherland"	"liberti"	"etern"
[4,]	"one"	"thee"	"flag"	"shall"	"live"
[5,]	"countri"	"bless"	"blood"	"hail"	"countri"
[6,]	"happi"	"free"	"arm"	"everi"	"free"
[7,]	"unit"	"sea"	"live"	"heart"	"sun"
[8,]	"uniti"	"thou"	"shall"	"defend"	"light"



## Examine beta: The per-topic-per-word probability

```
library(tidytext) # All three packages: "tidyverse"
library(ggplot2) # by Hadley Wickham library(tidyverse)

# A custom "tidier" for LDA output!

tidyTopics <- tidy(topic_model, matrix="beta")

# The %>% is "pipe" notation
anthem_top_topics <- tidyTopics %>%
  group_by(topic) %>%
  top_n(10, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)
```

```
> anthem_top_topics
# A tibble: 50 x 3
   topic term      beta
  <int> <chr>    <dbl>
1     1 us      0.131
2     1 peopl  0.0822
3     1 let    0.0795
4     1 one    0.0450
5     1 countri 0.0276
6     1 happi  0.0273
7     1 unit   0.0261
8     1 uniti  0.0247
9     1 aris   0.0213
10    1 great  0.0209
# ... with 40 more rows
```

# Use ggplot to visualize topics and terms



Beta is the probability that a given term appears in a particular topic. Higher probability terms "define" the topic best.

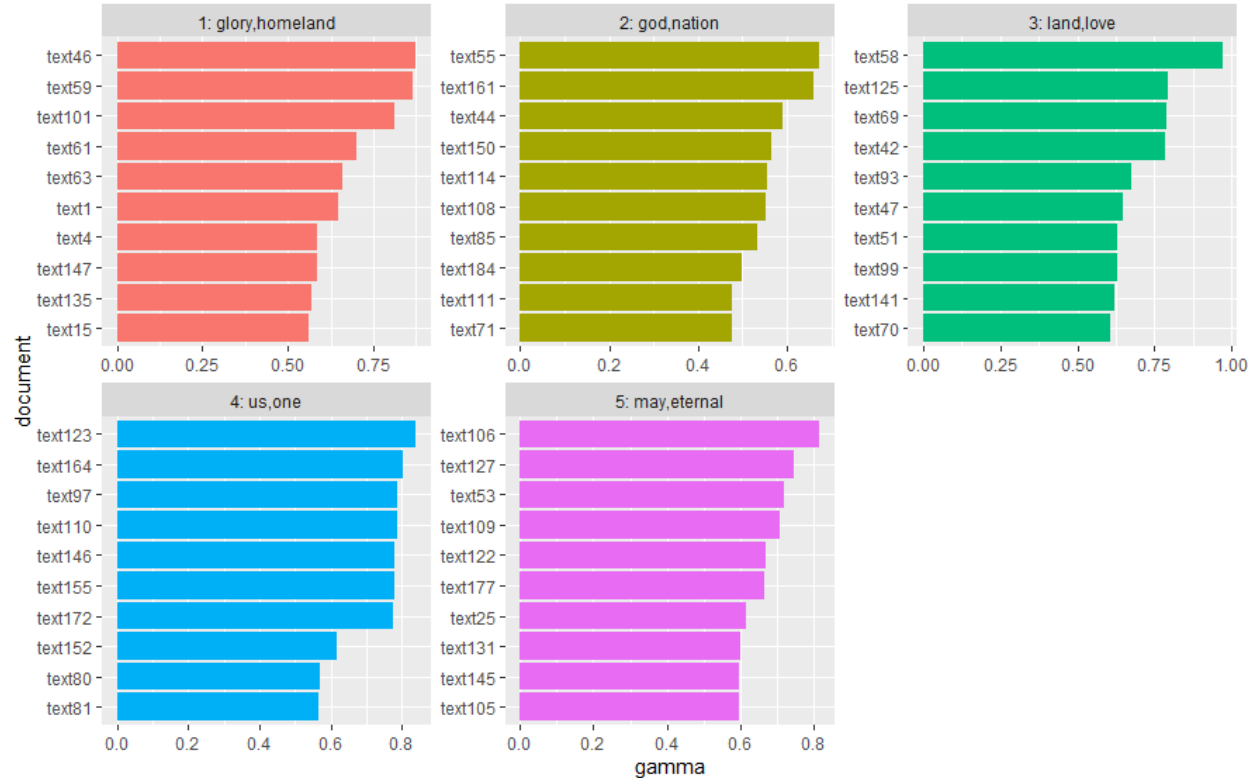
Note: your results will vary, but should reach to similar conclusions.

```
anthem_top_topics %>%  
  mutate(term = reorder(term, beta)) %>%  
  ggplot(aes(term, beta, fill = factor(topic)))  
  + geom_col(show.legend = FALSE) +  
  facet_wrap (~ topic, scales = "free") +  
  coord_flip()
```

## Try to name topics?

- Topic 1 top keywords: glory, homeland, fatherland
- My title for topic 1:  
*Glory to homeland*

# Use ggplot to visualize documents and topics



Gamma is the probability that a given topic appears in a particular document. Higher probability documents show the main topics.

```
anthem_top_topics %>%  
  mutate(document = reorder(document, gamma)) %>%  
  ggplot(aes(document, gamma, fill = factor(topic))) +  
  geom_col(show.legend = FALSE) +  
  facet_wrap(~ topic, scales = "free") +  
  coord_flip()
```

## Try to name topics?

- Topic 1 appears strongly in 3 anthems:
  - What countries are they?: Bolivia, Mexico, Iraq

*Of the homeland, the lofty name. In glorious splendour  
may we keep [it] And in its honour let's swear once  
again: To die before living as slaves!*

## Try to name topics?

- Topic 3?: *Land of love*
- Sri Lanka:

*Plenteous in prosperity, Thou, Beauteous in grace and love, Laden with grain and luscious fruit, And fragrant flowers of radiant hue, Giver of life and all good things, Our land of joy and victory, Receive our grateful praise sublime, we worship, worship Thee.*

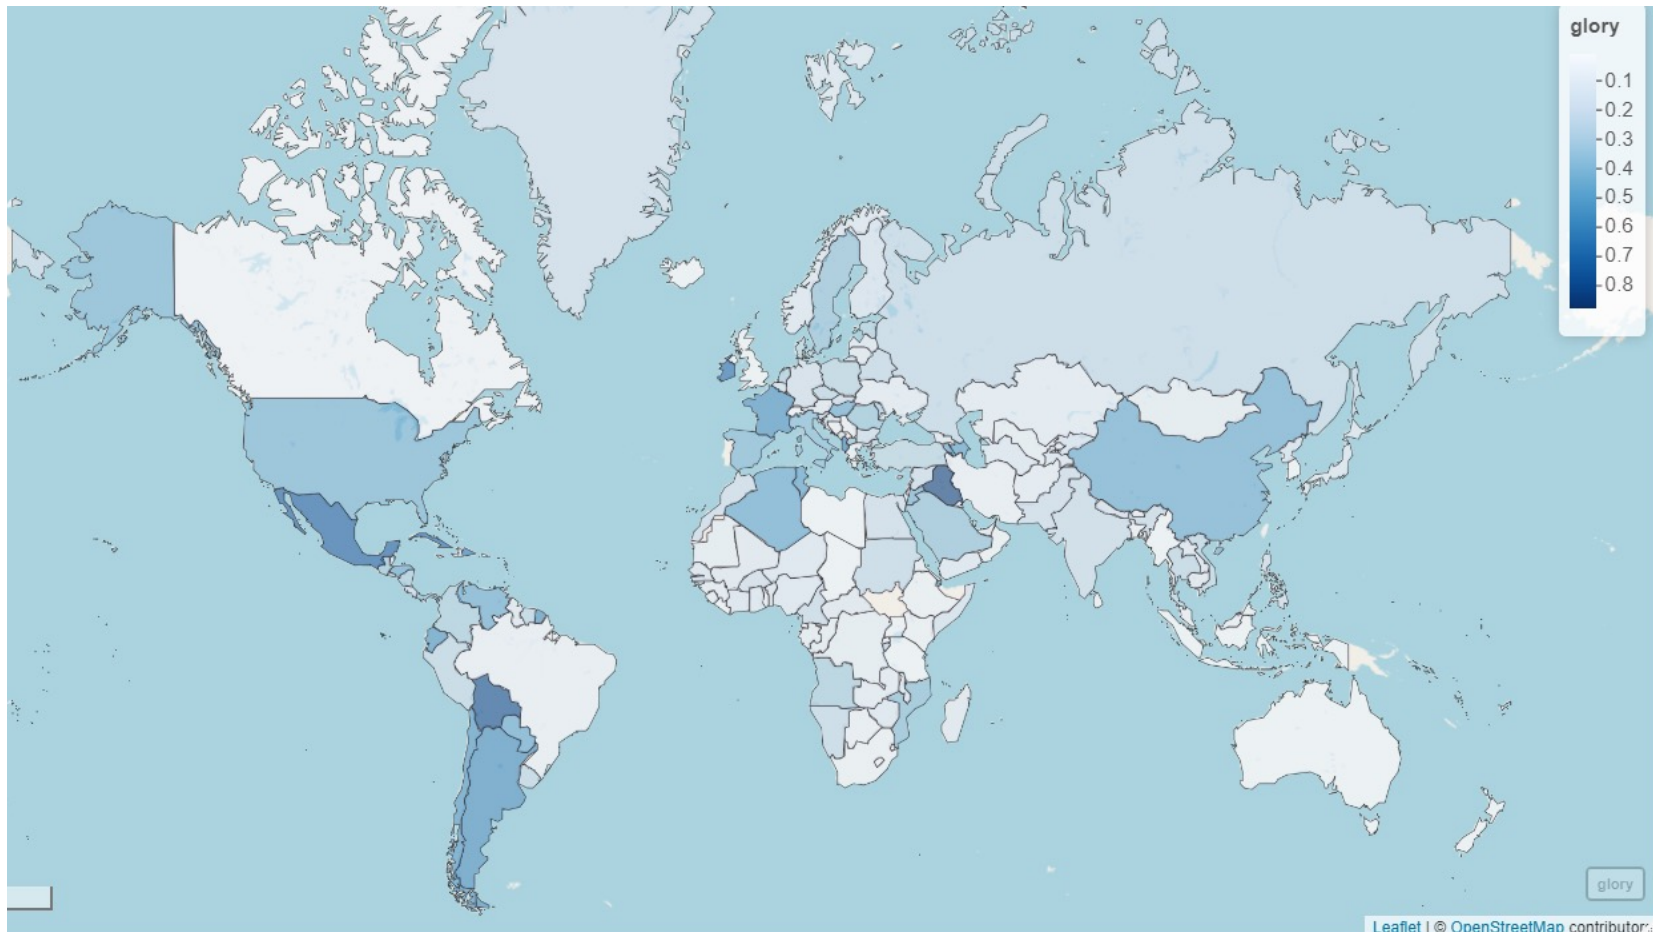
# Use ggplot to visualize documents and topics

- Note: For a given country, probabilities add up to 1:

```
> tidy_anthems[tidy_anthems$document=='text129',]  
# A tibble: 5 x 4  
  document topic  gamma topicname  
  <chr>    <int>  <dbl> <chr>  
1 text129      1 0.392 1: glory, homeland  
2 text129      2 0.254 2: god, nation  
3 text129      3 0.138 3: land, love  
4 text129      4 0.0874 4: us, one  
5 text129      5 0.128 5: may, eternal
```

# Try mapview()!!

```
library(mapview)  
mapview(out, zcol='gamma', col.regions = blues9,  
label=paste(out$CountryCode ,out$gamma), layer.name ='glory')
```





# Well.. with some cleaning

```
tidy_anthems <- tidy_anthems %>%
mutate(countryid=as.numeric(str_extract(tidy_anthems$document, "[0-9]+"))) # need the codes
anthems <- anthems %>%
  mutate(countryid = factor(rownames(anthems)))

df <- merge(tidy_anthems, anthems, on='countryid')
#####
#install.packages("rnaturalearth")
#install.packages("rnaturalearthdata")


library("rnaturalearth")
library("rnaturalearthdata")

world <- ne_countries(scale = "medium", returnclass = "sf")
world$CountryCode = world$su_a3
df$CountryCode = df$Alpha.3

out <- merge(df, world, on = 'CountryCode')
out <- out %>% filter(topic==1) %>% mutate(gamma=round(gamma,2))
# convert to spatial object
out <- st_as_sf(out)

library(mapview)
mapview(out, zcol='gamma', col.regions = blues9, label=paste(out$CountryCode
,out$gamma), layer.name = 'glory')
```

# Use Perplexity to Determine Number of Topics



```
library("topicmodels")
interview_topics <- convert(anthem_dtm, to = "topicmodels")
maxTopics <- 10 # Max number of topics we will allow to form
set.seed(1)

perList <- NULL # We are going to make a list of perplexity values
# Loop: Try every number of topics from 2 up to maxTopics
for (i in 2:maxTopics) {
  topic_model <- LDA(anthem_topics, method = "VEM", k=i) # Latent Dirichlet
  Allocation
  perList[i-1] <- perplexity(topic_model)
}

names(perList) <- as.character(2:maxTopics)
perList_df <- as.data.frame(perList)
library(ggplot2)
ggplot(data=perList_df, aes(x=as.numeric(rownames(perList_df)),
y=perList)) + geom_point() + geom_line(linetype='dashed') +
labs(x='num topics', y='perplexity')
```

This is an elbow plot!


Exercise: re-do the analysis with optimal k

# How Many Documents for Each Topic?

```
> top_anthems %>% group_by(topic) %>%  
  summarize(count=n(),av_gamma = mean(gamma))  
# A tibble: 5 × 3  
  topic count av_gamma  
  <int> <int>   <dbl>  
1     1    10    0.677  
2     2    10    0.608  
3     3    10    0.615  
4     4    10    0.685  
5     5    10    0.661
```

Ideally, the documents should be evenly distributed across topics.

```
top_anthems %>% group_by(topic) %>%  
  summarize(count=n(),av_gamma = mean(gamma)) %>%  
  ggplot(.,aes(x=topic,y=count)) + geom_col()
```



## **Exercise: Repeat analysis with nomination speech**

- Conduct topic analysis using the nominee text that you used for the sentiment analysis exercise
- Use paragraph numbers instead of descriptive titles for each of the paragraphs to use in place of anthems
- Generate LDA topic models and examine perplexity
- Visualize the topics and paragraphs and draw conclusions